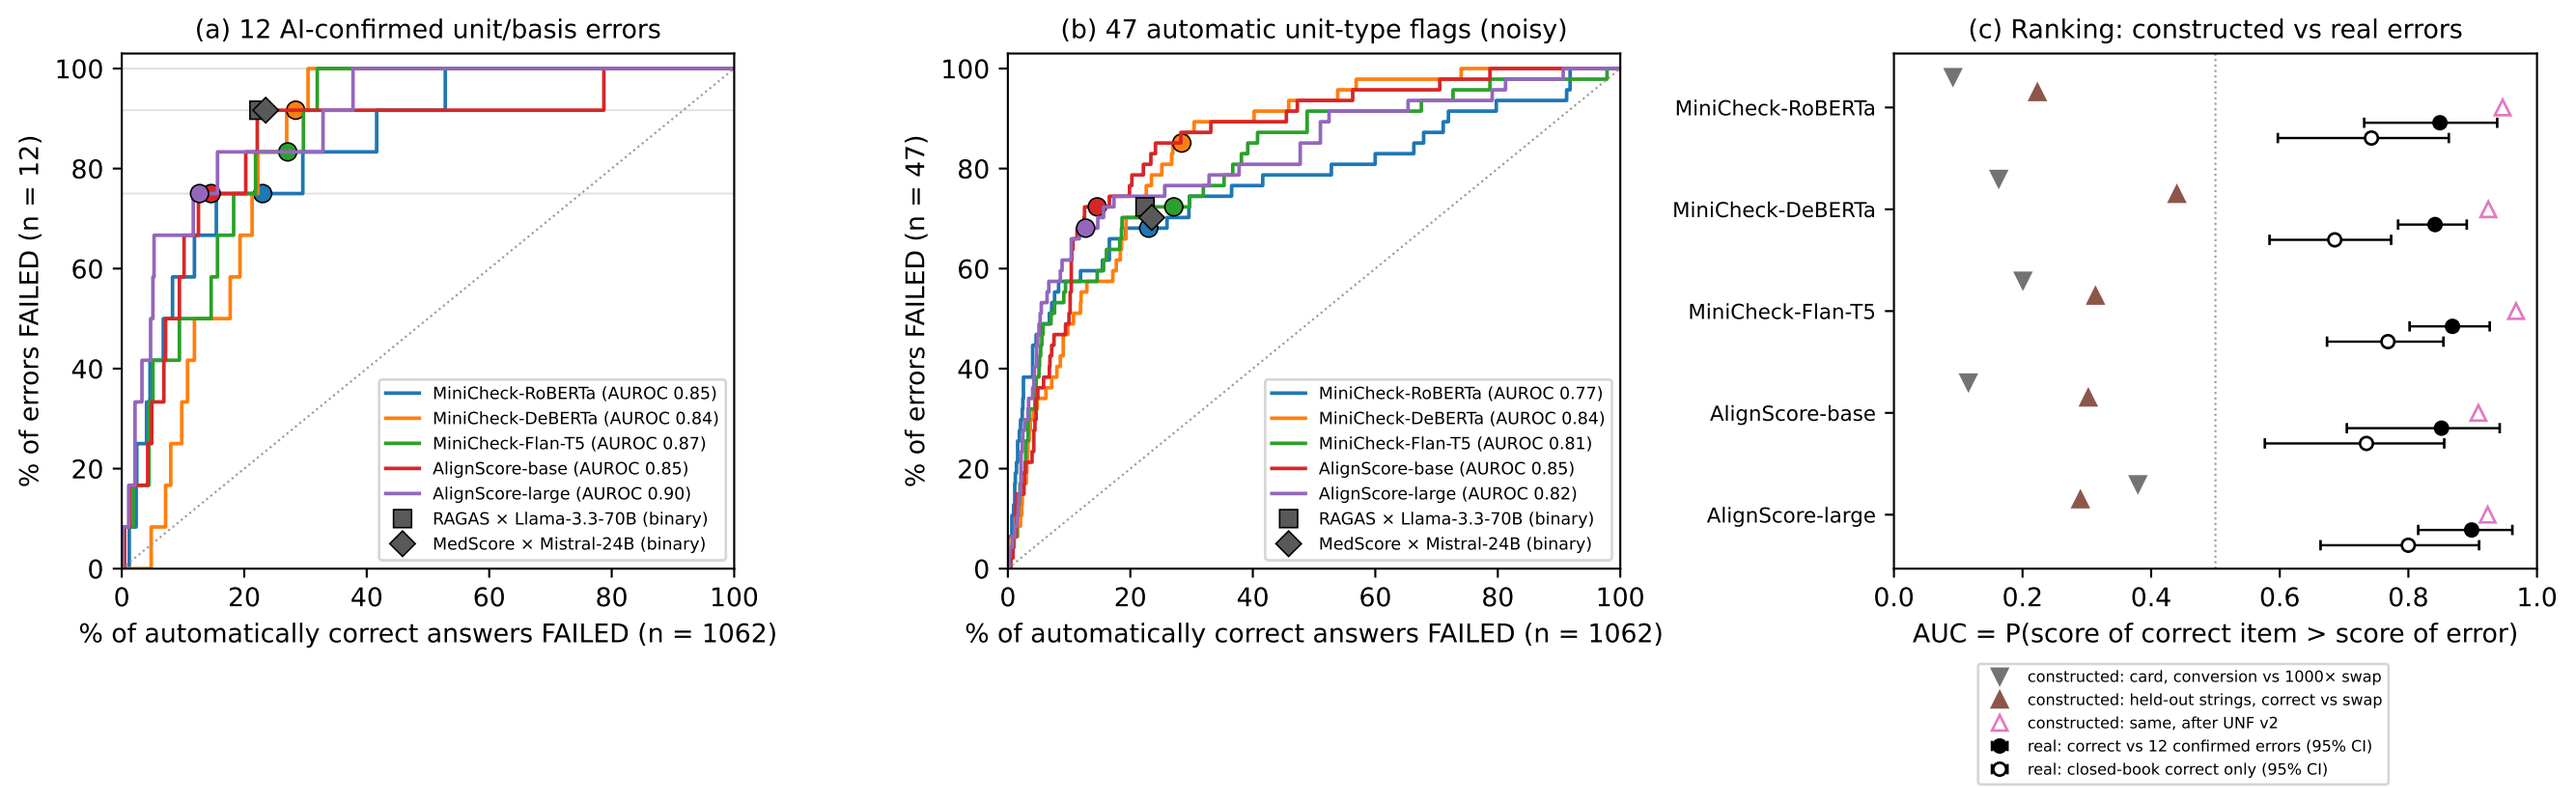
(a) 12 AI-confirmed unit/basis errors
% of errors FAILED (n = 12)
% of automatically correct answers FAILED (n = 1062)
MiniCheck-RoBERTa (AUROC 0.85)
MiniCheck-DeBERTa (AUROC 0.84)
MiniCheck-Flan-T5 (AUROC 0.87)
AlignScore-base (AUROC 0.85)
AlignScore-large (AUROC 0.90)
RAGAS × Llama-3.3-70B (binary)
MedScore × Mistral-24B (binary)
(b) 47 automatic unit-type flags (noisy)
% of errors FAILED (n = 47)
% of automatically correct answers FAILED (n = 1062)
MiniCheck-RoBERTa (AUROC 0.77)
MiniCheck-DeBERTa (AUROC 0.84)
MiniCheck-Flan-T5 (AUROC 0.81)
AlignScore-base (AUROC 0.85)
AlignScore-large (AUROC 0.82)
RAGAS × Llama-3.3-70B (binary)
MedScore × Mistral-24B (binary)
(c) Ranking: constructed vs real errors
MiniCheck-RoBERTa
MiniCheck-DeBERTa
MiniCheck-Flan-T5
AlignScore-base
AlignScore-large
AUC = P(score of correct item > score of error)
constructed: card, conversion vs 1000× swap
constructed: held-out strings, correct vs swap
constructed: same, after UNF v2
real: correct vs 12 confirmed errors (95% CI)
real: closed-book correct only (95% CI)